



UNIVERSITAT POLITÈCNICA DE CATALUNYA  
BARCELONATECH

Facultat d'Informàtica de Barcelona



## ADVANCED DATA STRUCTURES

COMPUTER SCIENCE DEPARTMENT

---

# Random Binary Search Trees: An empirical analysis

---

*Author:*

Alex Herrero

*Professor:*

Amalia Duch

March 8, 2025



## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Analysis of the Average Cost of Insertions</b>	<b>3</b>
2.1	Theoretical Study . . . . .	3
2.2	Experimental Study . . . . .	4
<b>3</b>	<b>Analysis of Internal Path Length</b>	<b>7</b>
3.1	Theoretical Study . . . . .	7
3.2	Experimental Study . . . . .	8
<b>4</b>	<b>Analysis of Interleaver Insertions and Deletions</b>	<b>10</b>
<b>A</b>	<b>More accurate solutions to the previous recurrences</b>	<b>11</b>

## List of Figures

1	Plot average number of nodes visited respect theoretical bound . . . . .	5
2	Plot average IPL respect the theoretical bound . . . . .	9
3	Plot experimental result divided by $n \log n$ . . . . .	9

## List of Tables

1	Experimental vs Theoretical results on the expected cost of insertion . . . . .	6
---	---------------------------------------------------------------------------------	---

## 1 Introduction

Let  $T$  be a binary tree with subtrees  $T_l$  and  $T_r$ . We say that  $T$  is a *binary search tree* (BST) if it is either an empty binary tree or it contains at least one element  $x$  as its root such that

- $T_l$  and  $T_r$  are also BSTs.
- $\forall y \in T_l, y < x$  and  $\forall z \in T_r, z > x$ .

Although it is well known that, in the worst case, a BST behaves like a linked list (with the height of the tree being  $\Theta(n)$ ), in this report, we focus on *random BSTs* of size  $n$ .

By *random BSTs*, we mean the following: Given a universe of keys  $U$  with  $|U| = n$ , we construct the BST by inserting each element of  $U$  exactly once, choosing the insertion order uniformly at random.

## 2 Analysis of the Average Cost of Insertions

### 2.1 Theoretical Study

Let us first analyze the expected cost of inserting an element  $u \in U$  into our BST  $T$ . For that, we will consider this cost as the cost of searching for  $u$  in our BST, which is valid since we can assume that, if  $u$  does not exist in  $T$ , our search terminates in any empty subtree with identical probability.

Let  $I_n$  be the expected cost of the insertion of a key  $x$  in a random BST of size  $n$ . Let, also, be  $I_{n,q}$  be the expected cost of the insertion of a key  $x$  in a random BST which root is the  $q$ -th smallest element. Then, the expected cost of  $I_n$  is

$$\begin{aligned}
 I_n &= \frac{1}{n} \sum_{q=1}^n I_{n,q} \\
 &= 1 + \frac{1}{n} \sum_{k=1}^n \left( \frac{1}{n} 0 + \frac{k-1}{n} I_{k-1} + \frac{n-k}{n} I_{n-k} \right) \\
 &= 1 + \frac{1}{n} \sum_{k=0}^{n-1} \left( \frac{k}{n} I_k + \frac{n-k-1}{n} I_{n-k-1} \right) \\
 &= 1 + \frac{1}{n^2} \sum_{k=0}^{n-1} (k I_k + (n-k-1) I_{n-k-1}) \\
 &= 1 + \frac{2}{n^2} \sum_{k=0}^{n-1} k I_k
 \end{aligned}$$

We can solve this recurrence using the continuous master theorem. The continuous master theorem solves recurrences of the form

$$F_n = t_n + \sum_{0 \leq j < n} w_{n,j} F_j$$

with  $t_n = \Theta(n^a (\log n)^b)$ . We proceed as follows:

- Determine the values of  $a$  and  $b$ : Since  $t_n = \Theta(1)$ , it is straightforward to see that  $a = b = 0$ .
- Provide a shape function for the weights  $w_{n,j}$ : We use the following trick to determine the shape function:

$$w(z) = \lim_{n \rightarrow \infty} n \cdot w_{n,z \cdot n} = n \cdot \frac{2zn}{n^2} = 2z.$$

- Determine the value of

$$\mathcal{H} = 1 - \int_0^1 w(z) z^a dz.$$

Substituting the values, we obtain:

$$\mathcal{H} = 1 - \int_0^1 2z dz = 1 - (1 - 0) = 0.$$

- Since  $\mathcal{H} = 0$ , we need to compute

$$\mathcal{H}' = -(b+1) \int_0^1 w(z) z^a \ln z dz.$$

Substituting the known values,

$$\mathcal{H}' = -1 \int_0^1 2z \ln z dz.$$

This integral can be solved using integration by parts. For the purpose of applying the theorem, we skip the detailed calculation, giving the result:

$$\mathcal{H}' = -\left(x^2 \ln x - \frac{x^2}{2}\right) \Big|_0^1 = \frac{1}{2}.$$

Since  $\mathcal{H} = 0$  and  $\mathcal{H}' \neq 0$ , we use the result

$$F_n = \frac{t_n}{\mathcal{H}'} \ln n + o(t_n \log n).$$

Substituting the values, we obtain

$$I_n = 2 \ln n + o(\log n).$$

Thus, the expected cost of an insertion into a random binary search tree is bounded by  $O(\log n)$ .

## 2.2 Experimental Study

Once we have theoretical results on the expected cost of an insertion in a random BST, we can provide experimental results to assess how closely they match the theoretical predictions. For this, we will conduct the following experiment:

1. We create a random BST of size  $n$  by generating  $n$  random keys in the interval  $[0, 1]$ .
2. After constructing the BST, we generate  $q = 2 \cdot n$  random numbers in the interval  $[0, 1]$ .
3. For each generated value, we perform a `find` operation in the BST, counting the number of nodes traversed during the operation.
4. We sum up the total number of nodes traversed across all  $q$  search operations and compute the average.
5. We repeat all previous steps with 20 different seeds and compute the final average.




Figure 1: Plot average number of nodes visited respect theoretical bound

6. We repeat the entire experiment for different values of  $n$ .

I conducted the experiment previously explained with values of  $n$  ranging from 10,000 to 200,000 in increments of 5,000, using seeds to generate random numbers from 1989 (in honor of the year of the first publication of the book *Introduction to Algorithms*, which was a great resource for refreshing my knowledge of BSTs and expected cost) to 2008. Figure 1 provides a plot of the values obtained from this experiment, as well as the theoretical bound derived using the continuous master theorem, indicating that the theoretical bound always bounds the experimental results. In fact, the shapes of both plots appear to be the same, suggesting that the only difference between the two functions is a constant factor. Indeed, by applying the continuous master theorem, we know that this difference is  $o(\log n)$ , but a more detailed analysis, presented in Appendix CITAR, shows that this difference is actually  $O(1)$ .

This constant factor is clearly visible in Table 1, where the column *Difference* (computed as the difference between the theoretical value obtained using Roura's theorem and the experimental value) consistently yields basically a constant difference.

$n$	Experimental	Theoretical	Difference
10000	17.6373	18.4207	0.78335
15000	18.4401	19.2316	0.79148
20000	19.0132	19.807	0.79379
25000	19.4589	20.2533	0.79437
30000	19.8232	20.6179	0.79466
35000	20.1327	20.9262	0.79355
40000	20.3994	21.1933	0.79389
45000	20.6318	21.4288	0.79707
50000	20.8435	21.6396	0.79606
55000	21.0332	21.8302	0.79698
60000	21.2081	22.0042	0.79606
65000	21.3695	22.1643	0.79484
70000	21.5175	22.3125	0.79501
75000	21.6545	22.4505	0.79602
80000	21.7826	22.5796	0.79696
85000	21.9042	22.7008	0.79657
90000	22.0176	22.8151	0.79753
95000	22.1266	22.9233	0.79666
100000	22.2288	23.0259	0.7971

$n$	Experimental	Theoretical	Difference
105000	22.3268	23.1234	0.79667
110000	22.4196	23.2165	0.79691
115000	22.5088	23.3054	0.79659
120000	22.5934	23.3905	0.79713
125000	22.6754	23.4721	0.7967
130000	22.7532	23.5506	0.79738
135000	22.8291	23.6261	0.79692
140000	22.901	23.6988	0.7978
145000	22.9706	23.769	0.79836
150000	23.0382	23.8368	0.79857
155000	23.1038	23.9024	0.79857
160000	23.1678	23.9659	0.79808
165000	23.2297	24.0274	0.79766
170000	23.2888	24.0871	0.79832
175000	23.3468	24.1451	0.79826
180000	23.4032	24.2014	0.79818
185000	23.4583	24.2562	0.79793
190000	23.5112	24.3096	0.79838
195000	23.5627	24.3615	0.79878
200000	23.6128	24.4121	0.79932

Table 1: Experimental vs Theoretical results on the expected cost of insertion

### 3 Analysis of Internal Path Length

#### 3.1 Theoretical Study

Let  $T$  be a BST of size  $n$  with root  $r$ . We define the *Internal Path Length* of  $T$  as the sum of all distances between every node of the tree and the root. More formally:

$$\text{IPL}(T) = \sum_{v \in V(T)} d(r, v).$$

This metric should not be surprising, as we can estimate the distance between the root and every node by averaging the IPL of such a tree. Before establishing this relationship, let us first derive a recurrence for the IPL of a tree, solve this recurrence, and later analyze the obtained relation.

Under the assumption of having a random BST, given  $n$  keys, we consider every permutation among the  $n!$  possible ones to be equally likely. Therefore, the sizes of the subtrees  $T_l$  and  $T_r$  will be completely random, provided that their combined sizes sum to  $n - 1$ . Consequently, in our recurrence, we must consider every possible size distribution of the subtrees as equally likely.

Moreover, given the IPL of  $T_l$  and  $T_r$ , since we introduce a new root at the top of the tree, every node will need to traverse one additional edge to reach the new root. This results in adding 1 to the previous distances with respect to their original root. As a consequence, we must add the number of edges in the tree to the sum of the previous IPLs of both subtrees.

Hence, creating the following recurrence:

$$\begin{aligned} \text{IPL}_n &= n - 1 + \frac{1}{n} \sum_{k=0}^{n-1} \text{IPL}_k + \text{IPL}_{n-1-k} \\ &= n - 1 + \frac{1}{n} \sum_{k=0}^{n-1} 2 \cdot \text{IPL}_k \\ &= n - 1 + \frac{2}{n} \sum_{k=0}^{n-1} \text{IPL}_k \end{aligned}$$

We can, again, use the continuous master theorem in order to solve this recurrence. Again we need to determinate the following values:

- Determine the values of  $a$  and  $b$ : Since  $t_n = \Theta(n)$ , it is straightforward to see that  $a = 1$  and  $b = 0$ .
- Provide a shape function for the weights  $w_{n,j}$ : We use the following trick to determine the shape function:

$$w(z) = \lim_{n \rightarrow \infty} n \cdot w_{n,z \cdot n} = n \cdot \frac{2}{n} = 2.$$

- Determine the value of

$$\mathcal{H} = 1 - \int_0^1 w(z) z^a dz.$$

Substituting the values, we obtain:

$$\mathcal{H} = 1 - \int_0^1 2z dz = 1 - (1 - 0) = 0$$

- Since  $\mathcal{H} < 0$  we are in the case

$$\mathcal{H}' = -(b+1) \int_0^1 w(z) z^a \ln z dz.$$

Substituting the known values,

$$\mathcal{H}' = -1 \int_0^1 2z \ln z dz.$$

This integral can be solved using integration by parts. For the purpose of applying the theorem, we skip the detailed calculation, giving the result:

$$\mathcal{H}' = -\left(x^2 \ln x - \frac{x^2}{2}\right) \Big|_0^1 = \frac{1}{2}.$$

Since  $\mathcal{H} = 0$  and  $\mathcal{H}' \neq 0$ , we use the result

$$F_n = \frac{t_n}{\mathcal{H}'} \ln n + o(t_n \log n).$$

Substituting the values, we obtain

$$I_n = 2(n-1) \ln n + o(n \log n) = 2n \ln n + o(n \log n)$$

Thus, the expected internal path length in a BST is bounded by  $O(n \log n)$ . Such recurrence can be used to solve the expected cost of doing an insertion in a BST as we can express that cost as:  $I_n = 1 + \frac{IPL_n}{n}$  (as we can have an estimation of  $I_n$  by averaging the internal path length over the size of the tree in order to get an estimation of the average distance between the root and a node) which would give a result of  $O(\log n)$  which was the expected bound we obtained in the previous section.

### 3.2 Experimental Study

Once we have theoretical results on the expected internal path length in a random BST, we can provide experimental results to assess how closely they match the theoretical predictions. For this, we will conduct the following experiment:

1. We create a random BST of size  $n$  by generating  $n$  random keys in the interval  $[0, 1]$ .
2. We calculate the internal path length of the random BST by doing a Breadth-first search algorithm
3. We repeat all previous steps with 20 different seeds and compute the final average.
4. We repeat the entire experiment for different values of  $n$ .

Again, I conducted the experiment with the same value of  $n$  and same values of *seeds* as before. Figure 2 provides, as before, a plot with the expected values we should obtain following the continuous master theorem result and the experimental result. This time we do not see a similar behavior between the two curves as we saw before: It seems that the theoretical bound increase faster than the experimental result, not at the same time as before.






Figure 2: Plot average IPL respect the theoretical bound

We know that there is a factor  $o(n \log n)$  (in fact, with a more detailed solution of the recurrence, see Appendix CITAR, we know that this  $o(n \log n)$  is  $O(n)$ ) that must be taken into account in our study. As a result, even with a high number of nodes in a random BST, this theoretical bound may still be influenced by the  $o(n \log n)$  factor until we reach a significantly larger number of nodes. To analyze this effect, we provide Figure 3, where we divide the experimental data by  $n \log n$  to observe when this  $o(n \log n)$  factor stabilizes. As seen in the figure, even for large input sizes (over 200,000 nodes), this factor still plays a significant role in our analysis.




Figure 3: Plot experimental result divided by  $n \log n$

So, that  $o(n \log n)$  factor must be the reason why it is difficult to predict a more accurate theoretical solution to the IPL: because this factor is very influential even in instances with more than 100,000 nodes, and we need to reach an even larger size of random BSTs to make this factor smaller than the actual bound.

## 4 Analysis of Interleaver Insertions and Deletions

Knuth in the volume 3 of his well-known book *The Art of Computer Programming* [Knu98] mentioned that, indeed, random insertions and deletions *destroys* the randomness of a BST. This fact was firstly observed by Gary Knott in 1972 and, later on, Eppinger [Epp83] showed expeerimentally that the path length, although at the beginning it decreases slightly, increases and reaches a stability after doing a quadratic number of deletions/insertions.



## A More accurate solutions to the previous recurrences

## References

- [Epp83] Jeffrey L Eppinger. An empirical study of insertion and deletion in binary search trees. *Communications of the ACM*, 26(9):663–669, 1983.
- [Knu98] Donald E Knuth. *The Art of Computer Programming: Sorting and Searching, volume 3*. Addison-Wesley Professional, 1998.